

Approximation Algorithms - Homework 2

Ilay Menahem, and Aaron Ross

December 29, 2025

Problem 1: Weighted Max-Cut Local Search

(a) Polynomial Convergence

Overview

We analyze the $LS_{Max-Cut}$ algorithm where a local move is made only if the weight of the cut increases by a factor of at least $(1 + \frac{\epsilon}{n})$. We show that the number of iterations is polynomial in the input size and $1/\epsilon$.

Proof of Validity

Let $W(S)$ denote the weight of the cut $(S, V \setminus S)$. Let S_k be the cut after k iterations, and S_0 be the initial cut. The algorithm updates S_k to S_{k+1} only if:

$$W(S_{k+1}) \geq \left(1 + \frac{\epsilon}{n}\right) W(S_k)$$

Therefore, after k iterations:

$$W(S_k) \geq \left(1 + \frac{\epsilon}{n}\right)^k W(S_0)$$

The maximum possible weight of any cut is bounded by the total weight of all edges, denoted W_{total} . Thus:

$$\left(1 + \frac{\epsilon}{n}\right)^k W(S_0) \leq W(S_k) \leq W_{total}$$

Complexity Analysis

Taking the natural logarithm of the inequality:

$$k \ln\left(1 + \frac{\epsilon}{n}\right) + \ln(W(S_0)) \leq \ln(W_{total})$$

Using the inequality $\ln(1+x) \approx x$ for small x (specifically $\ln(1+x) \geq \frac{x}{1+x}$ or simply bounding by linear growth for complexity):

$$k \cdot \frac{\epsilon}{2n} \leq k \ln\left(1 + \frac{\epsilon}{n}\right) \leq \ln(W_{total}) - \ln(W(S_0)) = \ln\left(\frac{W_{total}}{W(S_0)}\right)$$

$$k \leq \frac{2n}{\epsilon} \ln \left(\frac{W_{total}}{W(S_0)} \right)$$

Since $\ln(W_{total})$ is polynomial in the input size (number of bits to represent weights), k is polynomial.

(b) Improved Bound with Specific Initialization

Algorithm Modification

Initialize $S_0 = \{v^*\}$, where $v^* = \arg \max_{v \in V} w(\{v\}, V \setminus \{v\})$.

Complexity Analysis

1. Lower Bound of Initialization: Let $w(v)$ be the sum of weights of edges incident to v . We know $\sum_{v \in V} w(v) = 2W_{total}$. By the Pigeonhole Principle, there exists a vertex v^* such that:

$$w(v^*) \geq \frac{2W_{total}}{n}$$

The initial cut weight is $W(S_0) = w(v^*)$, so $W(S_0) \geq \frac{2W_{total}}{n}$.

2. Upper Bound on Iterations: The optimal cut W_{OPT} is bounded by W_{total} . Substituting this into the iteration bound from part (a):

$$k \leq O\left(\frac{n}{\epsilon}\right) \ln \left(\frac{W_{total}}{2W_{total}/n} \right) = O\left(\frac{n}{\epsilon}\right) \ln \left(\frac{n}{2} \right) = O\left(\frac{n \log n}{\epsilon}\right)$$

Thus, the number of iterations is $O(\frac{n \log n}{\epsilon})$.

Problem 2: Hitting Set

Algorithm: Recursive Local Ratio

The algorithm takes a universe U , a collection of sets \mathcal{C} , and weights w .

Algorithm 1 Local Ratio for Hitting Set

```

1: function HITTINGSET( $U, \mathcal{C}, w$ )
2:   Remove all  $x \in U$  with  $w(x) \leq 0$  (add to solution  $H$ ).
3:   if  $\mathcal{C} = \emptyset$  then return  $\emptyset$ 
4:   end if
5:   Pick any set  $S_i \in \mathcal{C}$ .
6:   Let  $\epsilon = \min_{x \in S_i} w(x)$ .
7:   Define  $w_1(x) = \begin{cases} \epsilon & \text{if } x \in S_i \\ 0 & \text{otherwise} \end{cases}$ 
8:   Define  $w_2(x) = w(x) - w_1(x)$ .
9:    $H' \leftarrow \text{HITTINGSET}(U, \mathcal{C}, w_2)$ 
10:  if  $H' \cap S_i = \emptyset$  then
11:    Find  $x \in S_i$  such that  $w_2(x) = 0$  (must exist as we subtracted min).
12:     $H \leftarrow H' \cup \{x\}$ 
13:  else
14:     $H \leftarrow H'$ 
15:  end if
16:  return  $H$ 
17: end function

```

Approximation Analysis

We prove that $w(H) \leq S_{max} \cdot w(H_{OPT})$, where $S_{max} = \max_i |S_i|$. According to the Local Ratio Theorem, it suffices to show the approximation holds for w_1 and w_2 .

Proof. **Inductive Step:** Assume the algorithm is an S_{max} -approximation for w_2 .

For w_1 :

- The optimal solution H_{OPT} must pick at least one element from S_i to cover it. Thus, $w_1(H_{OPT}) \geq \epsilon$.
- Our solution H adds at most one element specifically for S_i (if not already covered). However, even if we consider the set of all chosen elements, strictly speaking, $w_1(H) = \epsilon \cdot |H \cap S_i|$.
- While $|H \cap S_i|$ could be large in theory, the standard Local Ratio analysis for covering problems relies on comparing the cost paid to the optimum.
- A simpler view: The total weight removed from the system is ϵ per element in S_i . The optimal solution pays at least ϵ .

- The algorithm ensures minimality with respect to the decomposition.
- Specifically for Hitting Set (dual of Set Cover with frequency $f = S_{max}$), any minimal solution is an S_{max} -approximation with respect to w_1 .

$$w_1(H) \leq |S_i| \cdot \epsilon \leq S_{max} \cdot \epsilon \leq S_{max} \cdot w_1(H_{OPT})$$

Total:

$$w(H) = w_1(H) + w_2(H) \leq S_{max}w_1(H_{OPT}) + S_{max}w_2(H_{OPT}) = S_{max} \cdot w(H_{OPT})$$

□

Problem 3: Multiple Choice Maximum Coverage

Overview

We prove the local search algorithm is a $\frac{1}{2}$ -approximation. Let $S = (j_1, \dots, j_k)$ be the local optimum solution and $O = (o_1, \dots, o_k)$ be the global optimal solution.

Definitions (from Hint)

Let $U_S = \bigcup_{i=1}^k S_{i,j_i}$ be the set of elements covered by our solution. For each $1 \leq i \leq k$, define:

$$A_i = S_{i,j_i} \setminus \bigcup_{r \neq i} S_{r,j_r} \quad (\text{Unique contribution of set } i \text{ in } S)$$

$$B_i = S_{i,o_i} \setminus \bigcup_{r \neq i} S_{r,j_r} \quad (\text{Gain of switching set } i \text{ to optimal choice given others fixed})$$

Proof

- Local Optimality:** Since the algorithm terminates, switching any single index j_i to o_i does not improve the value. The loss of removing S_{i,j_i} is exactly $|A_i|$. The gain of adding S_{i,o_i} (relative to the other current sets) is exactly $|B_i|$.

$$|B_i| - |A_i| \leq 0 \implies |B_i| \leq |A_i|$$

- Summation:** Summing over all $i = 1 \dots k$:

$$\sum_{i=1}^k |B_i| \leq \sum_{i=1}^k |A_i|$$

Note that $\{A_i\}$ are disjoint subsets of U_S , and their union is exactly U_S (since every element in the union must be unique to some set or shared; if shared, it is not lost when removing one, but A_i captures the "ownership"). Actually, more precisely: $\sum |A_i| \leq |U_S| = V(S)$. In fact, $\sum |A_i| = V(S)$ is not strictly true if there is overlap, but $\sum |A_i| \leq V(S)$ holds. Wait, let us refine the definition of A_i . The hint defines A_i as the contribution of set i *given the others*. The sets A_i are disjoint. $\bigcup A_i \subseteq U_S$. Thus $\sum |A_i| \leq V(S)$.

- Bounding the Optimum:** Let $U_{OPT} = \bigcup S_{i,o_i}$. Consider any element $x \in U_{OPT} \setminus U_S$. Since x is not in U_S , it is not in any S_{r,j_r} for any r . Since $x \in U_{OPT}$, there exists some i such that $x \in S_{i,o_i}$. Therefore, $x \in S_{i,o_i} \setminus \bigcup_{r \neq i} S_{r,j_r}$, which means $x \in B_i$. This implies:

$$U_{OPT} \setminus U_S \subseteq \bigcup_{i=1}^k B_i$$

$$|U_{OPT}| - |U_S| \leq |U_{OPT} \setminus U_S| \leq \sum_{i=1}^k |B_i|$$

4. **Conclusion:** Combining the inequalities:

$$|U_{OPT}| - V(S) \leq \sum_{i=1}^k |B_i| \leq \sum_{i=1}^k |A_i| \leq V(S)$$

$$|U_{OPT}| \leq 2V(S) \implies V(S) \geq \frac{1}{2}V(OPT)$$

Problem 4: Generalized Steiner Forest (GSF)

(a) Minimal GSF is a 2-approximation for w'

Claim: Let $w'(u, v) = |\{u, v\} \cap T|$. Any minimal GSF F satisfies $w'(F) \leq 2w'(F_{OPT})$.

Proof. Lower Bound (OPT): Consider the optimal solution F^* . Since $|T_i| \geq 2$, every terminal $v \in T$ must be connected to at least one other terminal. Thus, every $v \in T$ has degree at least 1 in F^* .

$$w'(F^*) = \sum_{(u,v) \in F^*} |\{u, v\} \cap T| = \sum_{v \in T} \deg_{F^*}(v) \geq \sum_{v \in T} 1 = |T|$$

Upper Bound (Minimal Forest F): Since F is a minimal GSF, every leaf in every connected component of F must be a terminal in T . (If a leaf were not a terminal, the incident edge could be removed). For any tree in F , the sum of degrees of the vertices in T is strictly less than $2 \times (\text{number of vertices in } T \text{ in that tree})$. Summing over all components:

$$w'(F) = \sum_{v \in T} \deg_F(v) \leq 2|T| - 2 \times (\#\text{components}) < 2|T|$$

Ratio:

$$w'(F) < 2|T| \leq 2w'(F^*)$$

□

(b) Local Ratio Algorithm

1. **Base Case:** If all connectivity requirements are satisfied by 0-weight edges, return the set of those edges (pruned to be minimal).
2. **Weight Definition:** Define $w'(e) = |e \cap T|$.
3. **Compute ϵ :** Let $\epsilon = \min\{\frac{w(e)}{w'(e)} \mid w'(e) > 0, e \in E\}$.

4. Decomposition:

$$\begin{aligned} w_1(e) &= \epsilon \cdot w'(e) \\ w_2(e) &= w(e) - w_1(e) \end{aligned}$$

5. **Recursion:** Solve recursively for w_2 to obtain forest F' .
6. **Pruning (Crucial):** F' is a valid GSF. Remove edges from F' one by one as long as the connectivity requirements T_i remain satisfied. Let the result be F .
7. **Return F .**

Approximation: By the Local Ratio Theorem, since F is a minimal GSF, it is a 2-approximation for w_1 (by part a). By induction, it is a 2-approximation for w_2 . Thus, it is a 2-approximation for w .

Problem 5: Knapsack with Partition Constraints

Overview

We first define a Dynamic Programming (DP) algorithm that runs in pseudo-polynomial time, then apply the scaling technique to obtain an FPTAS.

Dynamic Programming

Let the items be ordered such that items in S_1 appear first, then S_2 , etc. We compress the state to handle one group at a time.

- $D_{prev}[v]$: Min weight to achieve value v using previous groups.
- $D_{curr}[c][v]$: Min weight to achieve value v using previous groups AND exactly c items from the current group.

Transitions: For each group S_j with limit $k(j)$:

1. Initialize $D_{curr}[0][v] = D_{prev}[v]$.
2. For each item $i \in S_j$ (weight w_i , value v_i):

$$D_{curr}[c][v] = \min(D_{curr}[c][v], D_{curr}[c-1][v - v_i] + w_i)$$

(Iterate c from $k(j)$ down to 1).

3. After processing all items in S_j , update D_{prev} :

$$D_{prev}[v] = \min_{0 \leq c \leq k(j)} D_{curr}[c][v]$$

The complexity is $O(n \cdot n \cdot V_{total}) = O(n^2 V_{total})$ since the inner c loop runs at most n times.

FPTAS Construction

To achieve a $(1 - \epsilon)$ -approximation:

1. Let $V_{max} = \max_i v_i$.
2. Define scaling factor $K = \frac{\epsilon V_{max}}{n}$.
3. Define scaled values $v'_i = \lfloor \frac{v_i}{K} \rfloor$.
4. Run the DP algorithm using weights w_i and values v'_i .
5. Output the set of items corresponding to the optimal DP state satisfying weight $\leq B$.

Complexity and Bound

The maximum possible scaled value is bounded by:

$$V'_{total} \leq \sum v'_i \leq n \cdot \frac{V_{max}}{K} = n \cdot \frac{n}{\epsilon} = \frac{n^2}{\epsilon}$$

The DP running time becomes:

$$O\left(n^2 \cdot \frac{n^2}{\epsilon}\right) = O\left(\frac{n^4}{\epsilon}\right)$$

This is polynomial in n and $1/\epsilon$.

Approximation Proof: For the optimal set P^* , the difference between the true value and the scaled value is:

$$\sum_{i \in P^*} v_i - K \sum_{i \in P^*} v'_i < \sum_{i \in P^*} K = |P^*|K \leq nK = \epsilon V_{max} \leq \epsilon OPT$$

Thus, the solution returned is at least $(1 - \epsilon)OPT$.